

Machine Learning & Data-Mining : Research Project

Implementation and experiments

Predictions of diseases using decision trees

References for this project :

- Chronic kidney disease diagnosis using decision trees algorithms (2021) (<https://bmcnephrol.biomedcentral.com/articles/10.1186/s12882-021-02474-z>)
- Simple Prediction of Type 2 Diabetes Mellitus via Decision Tree Modeling (2017) (<https://brieflands.com/articles/ircrj-10657.pdf>)
- Decision tree model in the diagnosis of breast cancer (2017) (<https://ieeexplore.ieee.org/document/8789297>)
- Early Prediction of Heart Disease Using Decision Tree Algorithm (2017) (https://www.researchgate.net/profile/Safish-Mary/publication/315023624_Early_Prediction_of_Heart_Disease_Using_Decision_Tree_Algorithm/links/58c84b57aca2723ab16eba60/Early-Prediction-of-Heart-Disease-Using-Decision-Tree-Algorithm.pdf)

I will get help with the code of the first assignment on decision trees for my implementation.

I am re-implementing what I coded with another dataset. I used the dataset « Breast cancer Wisconsin » : <https://www.kaggle.com/datasets/uciml/breast-cancer-wisconsin-data>

Let's highlight the key components of the code and their functionalities:

- Libraries :

```
6 import pandas as pd
7 from typing import Union
8 import matplotlib.pyplot as plt
9 import numpy as np
10 from sklearn.tree import DecisionTreeClassifier, plot_tree
11 from sklearn.metrics import accuracy_score
```

- Data Loading and Preprocessing:

The `load_data` function loads the Breast Cancer dataset from a CSV file. This function drops unnecessary columns from the dataset and features and targets are extracted.

```
201 def load_data():
202     BreastCancerDataset = pd.read_csv('BreastCancerDataset.csv').drop(['Unnamed: 32'], axis=1)
203     BreastCancerFeatures = BreastCancerDataset.drop(['id', 'diagnosis'], axis=1).values
204     BreastCancerTargets = BreastCancerDataset['diagnosis'].values
205     BreastCancerClasses = BreastCancerDataset['diagnosis'].unique()
206     return BreastCancerFeatures, BreastCancerTargets, BreastCancerClasses
```

- Prior probability function

The prior function calculates the prior probability of each class type in a classification problem. It represents the distribution of classes in the dataset before considering any features.

```

35 def prior(targets: np.ndarray, classes: list) -> np.ndarray:
36     """
37     Calculate the prior probability of each class type
38     given a list of all targets and all class types
39     """
40     class_counts = np.zeros(len(classes))
41
42     for target in targets:
43         class_counts[target] += 1
44
45     total_samples = len(targets)
46     class_probabilities = class_counts / total_samples
47
48     return class_probabilities

```

- Splitting Data:

The `split_train_test` function shuffles and splits the dataset into training and testing sets based on a specified ratio. I am using the same `split_train_test` function from the tools of the assignment.

```

13 def split_train_test(features: np.ndarray, targets: np.ndarray,
14                      train_ratio: float=0.8) -> Union[tuple, tuple]:
15     """
16     Shuffle the features and targets in unison and return
17     two tuples of datasets, first being the training set,
18     where the number of items in the training set is according
19     to the given train_ratio
20     """
21     np.random.seed(888)
22     p = np.random.permutation(features.shape[0])
23     features = features[p]
24     targets = targets[p]
25
26     split_index = int(features.shape[0] * train_ratio)
27
28     train_features, train_targets = features[0:split_index, :], \
29     targets[0:split_index]
30     test_features, test_targets = features[split_index:-1, :], \
31     targets[split_index:-1]
32
33     return (train_features, train_targets), (test_features, test_targets)
34

```

- Decision Tree Functions:

- `split_data` Function:

This function splits the dataset and targets into two separate datasets based on a given feature and threshold.

```

51 def split_data(
52     features: np.ndarray,
53     targets: np.ndarray,
54     split_feature_index: int,
55     theta: float
56 ) -> Union[tuple, tuple]:
57     """
58     Split a dataset and targets into two separate datasets
59     where data with split_feature < theta goes to 1 otherwise 2
60     """
61     split1 = features[:, split_feature_index] < theta
62     split2 = features[:, split_feature_index] >= theta
63
64     features_1 = features[split1]
65     targets_1 = targets[split1]
66
67     features_2 = features[split2]
68     targets_2 = targets[split2]
69
70     return (features_1, targets_1), (features_2, targets_2)

```

➤ gini_impurity Function:

This function calculates the Gini impurity for a set of targets.

```

73 def gini_impurity(targets: np.ndarray, classes: list) -> float:
74     '''
75     Calculate:
76      $i(S_k) = 1/2 * (1 - \sum_i P\{C_i\}^2)$ 
77     '''
78
79     square = prior(targets, classes)
80     for n in range(len(square)): square[n] = square[n]**2
81
82     return 1/2 * (1 - sum(square))

```

➤ weighted_impurity Function:

This function calculates the weighted sum of Gini impurities for two branches.

```

85 def weighted_impurity(
86     t1: np.ndarray,
87     t2: np.ndarray,
88     classes: list
89 ) -> float:
90     '''
91     Given targets of two branches, return the weighted
92     sum of gini branch impurities
93     '''
94
95     g1 = gini_impurity(t1, classes)
96     g2 = gini_impurity(t2, classes)
97     n = t1.shape[0] + t2.shape[0]
98
99     i = (t1.shape[0] * g1 + t2.shape[0] * g2) / n
100     return i

```

➤ total_gini_impurity Function:

This function calculates the total Gini impurity which returns the weighted impurity given the dataset and threshold to split on.

```

102 def total_gini_impurity(
103     features: np.ndarray,
104     targets: np.ndarray,
105     classes: list,
106     split_feature_index: int,
107     theta: float
108 ) -> float:
109     '''
110     Calculate the gini impurity for a split on split_feature_index
111     for a given dataset of features and targets.
112     '''
113
114     (features_1, targets_1), (features_2, targets_2) = split_data(features, targets, split_feature_index, theta)
115     weight = weighted_impurity(targets_1, targets_2, classes)
116
117     return weight

```

➤ brute_best_split Function:

This function finds the best split for the given data by iterating over feature dimensions and thresholds.

```

120 def brute_best_split(
121     features: np.ndarray,
122     targets: np.ndarray,
123     classes: list,
124     num_tries: int
125 ) -> Union[float, int, float]:
126     """
127     Find the best split for the given data. Test splitting
128     on each feature dimension num_tries times.
129
130     Return the lowest gini impurity, the feature dimension and
131     the threshold
132     """
133
134     best_gini, best_dim, best_theta = float("inf"), None, None
135
136     # iterate feature dimensions
137     for i in range(features.shape[1]):
138         features_i = features[:,i]
139
140         # create the thresholds
141         thetas = np.linspace(features_i.min(), features_i.max(), num_tries+2)[1:-1]
142
143         # iterate thresholds
144         for theta in thetas:
145             gini = total_gini_impurity(features, targets, classes, i, theta)
146             if best_gini > gini:
147                 best_gini = gini
148                 best_dim = i
149                 best_theta = theta
150
151     return best_gini, best_dim, best_theta

```

- Decision Tree Trainer Class:

This class handles the training and evaluation of the Decision Tree model of the breast cancer winsconsin dataset. It utilizes the DecisionTreeClassifier from scikit-learn and provides methods for training, accuracy calculation, plotting, making predictions, and generating a confusion matrix.

```


154 class BreastCancerTreeTrainer:
155     def __init__(
156         self,
157         features: np.ndarray,
158         targets: np.ndarray,
159         classes: list = ['M', 'B'],
160         train_ratio: float = 0.8
161     ):
162         """
163         train_ratio: The ratio of the Breast Cancer dataset that will
164         be dedicated to training.
165         """
166         (self.train_features, self.train_targets), \
167         (self.test_features, self.test_targets) = \
168             split_train_test(features, targets, train_ratio)
169
170         self.classes = classes
171         self.tree = DecisionTreeClassifier()
172
173     def train(self):
174         return self.tree.fit(self.train_features, self.train_targets)
175
176     def accuracy(self):
177         return accuracy_score(self.test_targets, self.tree.predict(self.test_features))
178
179     def plot(self):
180         plt.figure(figsize=(10,10), dpi=300)
181         plot_tree(self.tree, filled=True, class_names=self.classes, rounded=True)
182         plt.show()
183
184     def guess(self):
185         return self.tree.predict(self.test_features)
186
187     def confusion_matrix(self):
188         predictions = self.guess()
189         num_classes = len(self.classes)
190         cm = np.zeros((num_classes, num_classes), dtype=int)
191
192         for true_label, pred_label in zip(self.test_targets, predictions):
193             cm[true_label][pred_label] += 1
194
195         return cm
196

```

- Main Script:

In the main, I load data using the `load_data` function. I instantiate a `BreastCancerTreeTrainer` object and train a Decision Tree model in order to do prediction and classify if the patient's tumor is benign or malignant. If it is malignant, it means that the patient has breast cancer.

For the moment, I have this decision tree while taking into consideration all the features of the Breast Cancer Dataset:



The accuracy achieved by my implementation is 95.58%, slightly surpassing the 94.3% reported in the referenced article (<https://ieeexplore.ieee.org/document/9442043/>). Despite this high accuracy, it's crucial to acknowledge the presence of bugs in my code, potentially leading to inaccurate results.

One identified issue lies in the prior function, which requires modification. The breast cancer dataset's classes, denoted as 'B' for benign and 'M' for malignant, are of string type. This contrasts with the integer type classes in the iris dataset. Consequently, the prior function needs adjustments to handle string-type classes.


Additionally, the confusion matrix is currently non-functional. Addressing this issue is essential for a comprehensive evaluation of the model's performance.

Aside of this first plot, I tried to see what features of the dataset have the biggest impact on results in order to reduce the number of features (more than 30). By retaining only features that are not correlated, the model can improve in interpretability.


This function helped me in obtaining a correlation map of the Breast Cancer dataset's features.

```
198 import seaborn as sns
199
200 def feature_correlation_map(dataframe):
201     # Calculate the correlation matrix
202     corr_matrix = dataframe.corr()
203
204     # Create a mask for the upper triangle
205     mask = np.triu(np.ones_like(corr_matrix, dtype=bool))
206
207     # Set up the matplotlib figure
208     plt.figure(figsize=(12, 10))
209
210     # Draw the heatmap with the mask
211     sns.heatmap(corr_matrix, mask=mask, cmap='coolwarm', vmax=1, center=0,
212               square=True, linewidths=.5, cbar_kws={"shrink": 0.5})
213
214     plt.show()
```

I applied the feature_correlation_map to my Breast Cancer Dataset and I obtained this map :



After retaining only the features that has a correlation higher than the threshold of 0.5, we obtain a new correlation map :



Radius_mean, texture_mean, smoothness_mean, texture_se, smoothness_se and symmetry_se are the features which are not correlated.

The next step of my implementation will be to correct bugs that I mention earlier and to understand more deeply how to interpret each decision tree's plot.